



# GPU BOOTCAMP

KTH September 2019

Paul Graham : [pgraham@nvidia.com](mailto:pgraham@nvidia.com)

# INTRODUCTION TO GPU COMPUTING

## What to expect?

- Broad view on GPU Stack
- Fundamentals of GPU Architecture
- Ways to GPU Computing
- Good starting point

# FULL STACK OPTIMIZATION


Progress Of Stack In 6 Years

2013

cuBLAS: 5.0
cuFFT: 5.0
cuRAND: 5.0
cuSPARSE: 5.0
NPP: 5.0
Thrust: 1.5.3
CUDA: 5.0
Resource Mgr: r304
Base OS: CentOS 6.2



Accelerated Server  
With Fermi



2019

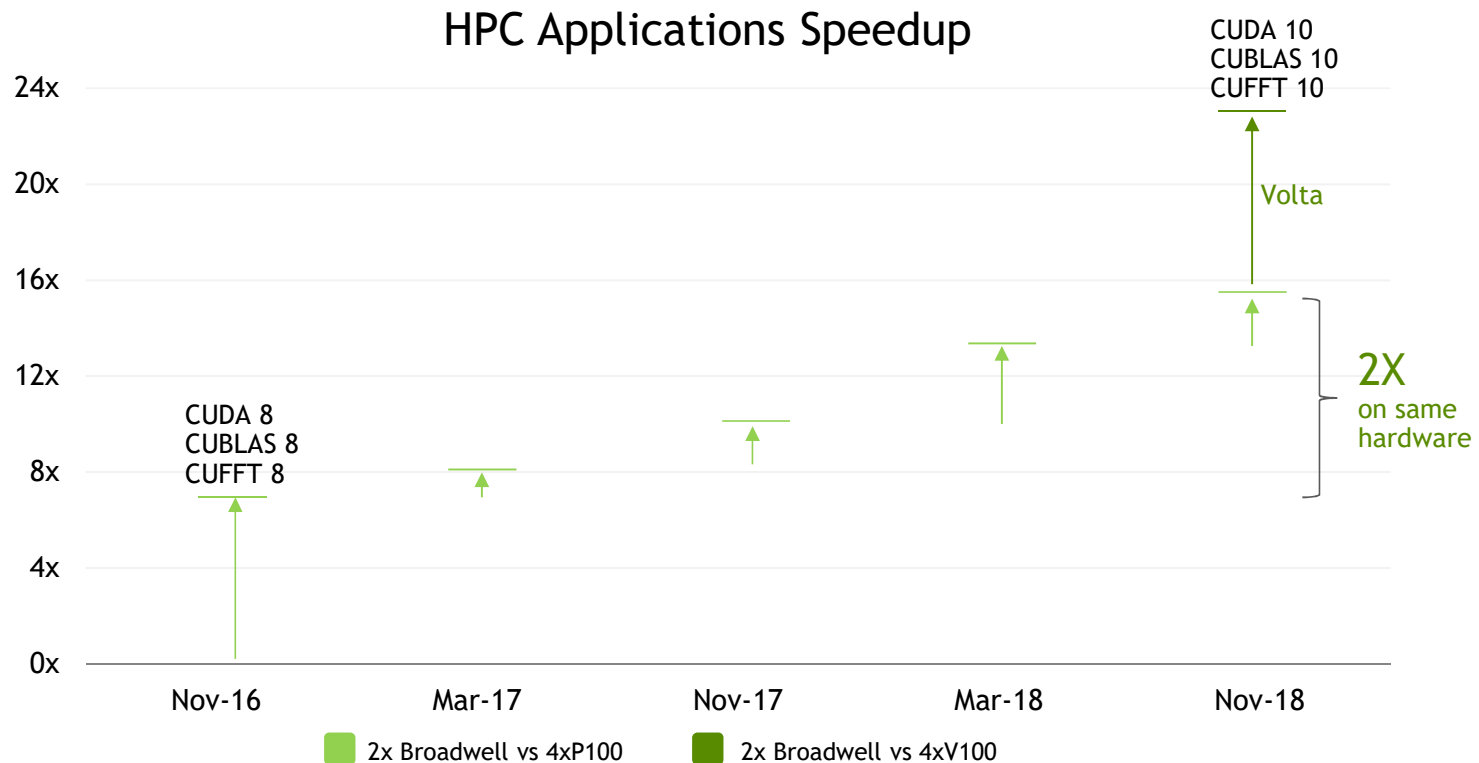
cuBLAS: 10.0
cuFFT: 10.0
cuRAND: 10.0
cuSOLVER: 10.0
cuSPARSE: 10.0
NPP: 10.0
Thrust: 1.9.0
CUDA: 10.0
Resource Mgr: r384
Base OS: Ubuntu 16.04



Accelerated Server  
with Volta

# ACCELERATED COMPUTING IS FULL-STACK OPTIMIZATION

2X More Performance with Software Optimizations Alone



HPC Apps: AMBER, Chroma, GROMACS, GTC, LAMMPS, MILC, NAMD, QE, RTM, SPECfem3D, VASP

# NVIDIA UNIVERSAL ACCELERATION PLATFORM

Single Platform Drives Utilization and Productivity


## CUSTOMER USECASES



Speech




Translate



Recommender

CONSUMER INTERNET



Healthcare



Manufacturing



Finance


INDUSTRIAL APPLICATIONS



Molecular Simulations



Weather Forecasting



Seismic Mapping

SCIENTIFIC APPLICATIONS

## APPS & FRAMEWORKS



## NVIDIA SDK & LIBRARIES

### MACHINE LEARNING/ ANALYTICS

cuDF

cuML

cuGRAPH

### DEEP LEARNING

cuDNN

cuBLAS

CUTLASS

NCCL

TensorRT

### HPC

CuBLAS

CuFFT

OpenACC

CUDA

## TESLA GPUs & SYSTEMS



TESLA GPU



VIRTUAL GPU



NVIDIA DGX FAMILY



NVIDIA HGX

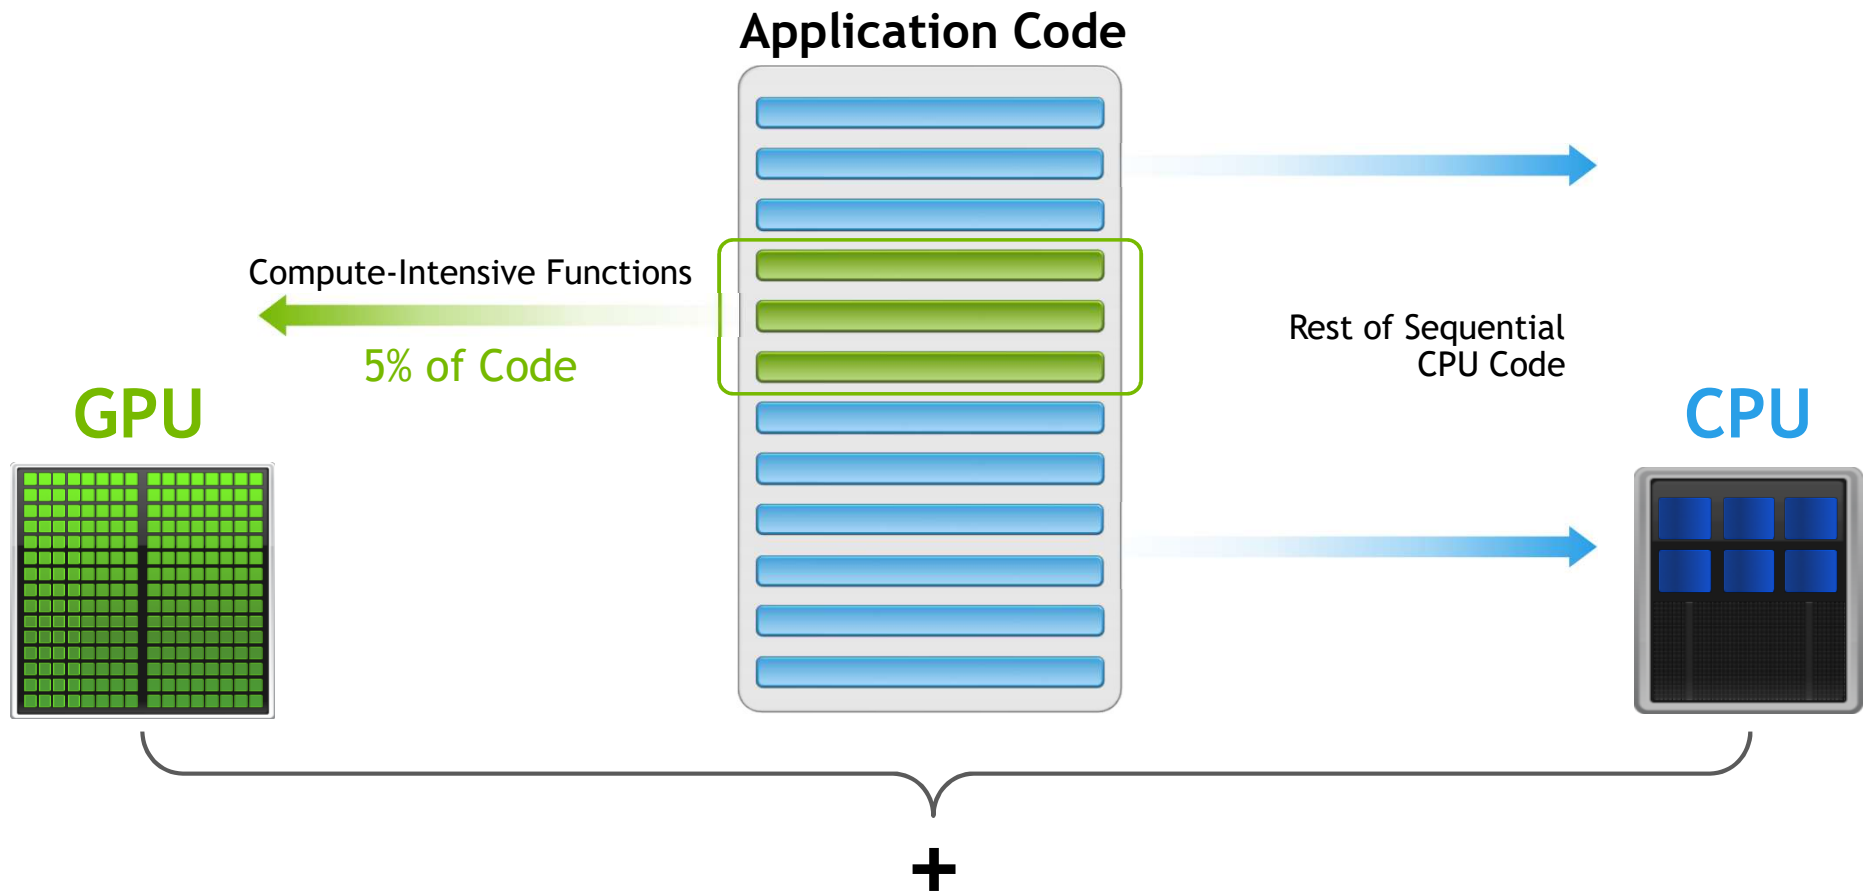


SYSTEM OEM



CLOUD

# HOW GPU ACCELERATION WORKS






# ACCELERATED COMPUTING


## CPU

Optimized for  
Serial Tasks



## GPU Accelerator


Optimized for  
Parallel Tasks



# CPU IS A LATENCY REDUCING ARCHITECTURE

## CPU

Optimized for  
Serial Tasks




## CPU Strengths

- Very large main memory
- Very fast clock speeds
- Latency optimized via large caches
- Small number of threads can run very quickly

## CPU Weaknesses

- Relatively low memory bandwidth
- Cache misses very costly
- Low performance/watt





# GPU IS ALL ABOUT HIDING LATENCY

## GPU Strengths




- High bandwidth main memory
- Significantly more compute resources
- Latency tolerant via parallelism
- High throughput
- High performance/watt

## GPU Weaknesses

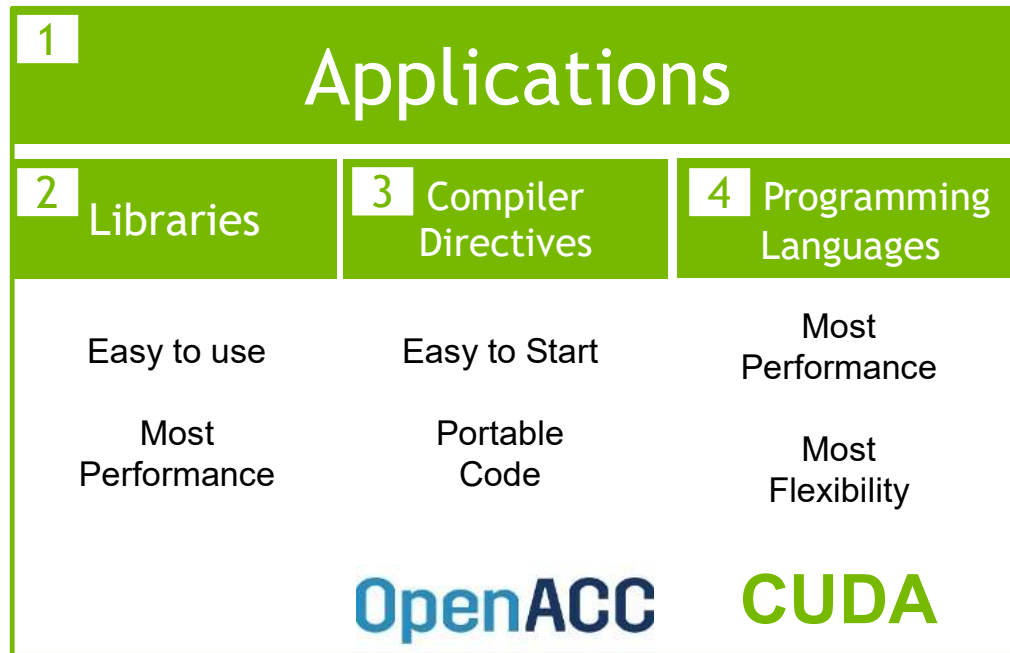
- Relatively low memory capacity
- Low per-thread performance

## GPU Accelerator

Optimized for  
Parallel Tasks



# HOW TO START WITH GPUS



1. Review available GPU-accelerated applications
2. Check for GPU-Accelerated applications and libraries
3. Add OpenACC Directives for quick acceleration results and portability
4. Dive into CUDA for highest performance and flexibility

The background is a dark, almost black, field filled with a complex network of thin, glowing green lines. These lines intersect at various points, creating a web-like structure. At several of these intersection points, there are small, bright green circular dots. Additionally, there are a few larger, fainter blue circular shapes scattered across the background, particularly in the upper right quadrant. The overall effect is one of a dynamic, interconnected system, possibly representing a network or a simulation.

# **GPU COMPUTING PLATFORM FOR HPC SIMULATION**

# GPU-ACCELERATED APPLICATIONS


## 620 Applications Across Domains

- ▶ Life Sciences
- ▶ Manufacturing
- ▶ Physics
- ▶ Oil & Gas
- ▶ Climate & Weather
- ▶ Media & Entertainment
- ▶ Deep Learning
- ▶ Federal & Defense
- ▶ Data Science & Analytics
- ▶ Safety & Security
- ▶ Computational Finance
- ▶ Tool & Management


# GPU ACCELERATED LIBRARIES

“Drop-in” Acceleration for Your Applications


## DEEP LEARNING




## SIGNAL, IMAGE & VIDEO



## LINEAR ALGEBRA



## PARALLEL ALGORITHMS



More libraries: <https://developer.nvidia.com/gpu-accelerated-libraries>

# WHAT IS OPENACC

Programming Model for an Easy Onramp to GPUs

Directives-based  
programming model for  
**parallel  
computing**

Add Simple Compiler Directive

```
main()
{
  <serial code>
  #pragma acc kernels
  {
    <parallel code>
  }
}
```

Simple

Designed for  
**performance  
portability** on  
CPUs and GPUs

Powerful & Portable

Read more at [www.openacc.org/about](http://www.openacc.org/about)

OpenACC is an open specification developed by OpenACC.org consortium



# SINGLE PRECISION ALPHA X PLUS Y (SAXPY)

GPU SAXPY in multiple languages and libraries

Part of Basic Linear Algebra Subroutines (BLAS) Library

$$z = \alpha x + y$$

$x, y, z$  : vector

$\alpha$  : scalar

## 1

# SAXPY: OPENACC COMPILER DIRECTIVES

## Parallel C Code

```
void saxpy(int n,
           float a,
           float *x,
           float *y)
{
    #pragma acc kernels
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

...
// Perform SAXPY on 1M elements
saxpy(1<<20, 2.0, x, y);
...
```

## Parallel Fortran Code

```
subroutine saxpy(n, a, x, y)
    real :: x(:), y(:), a
    integer :: n, i
    !$acc kernels
    do i=1,n
        y(i) = a*x(i)+y(i)
    enddo
    !$acc end kernels
end subroutine saxpy

...
! Perform SAXPY on 1M elements
call saxpy(2**20, 2.0, x_d, y_d)
...
```

## 2

# SAXPY: CUBLAS LIBRARY

## Serial BLAS Code

```
int N = 1<<20;

...

// Use your choice of blas library

// Perform SAXPY on 1M elements
blas_saxpy(N, 2.0, x, 1, y, 1);
```

## Parallel cuBLAS Code

```
int N = 1<<20;

cublasInit();
cublasSetVector(N, sizeof(x[0]), x, 1, d_x, 1);
cublasSetVector(N, sizeof(y[0]), y, 1, d_y, 1);

// Perform SAXPY on 1M elements
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);

cublasGetVector(N, sizeof(y[0]), d_y, 1, y, 1);

cublasShutdown();
```

You can also call cuBLAS from Fortran, C++, Python, and other languages:

<http://developer.nvidia.com/cublas>

## 3

## SAXPY: CUDA C

## Standard C

```

void saxpy(int n, float a,
          float *x, float *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}

int N = 1<<20;

// Perform SAXPY on 1M elements
saxpy(N, 2.0, x, y);

```

## Parallel C

```

__global__
void saxpy(int n, float a,
          float *x, float *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n) y[i] = a*x[i] + y[i];
}

int N = 1<<20;
cudaMemcpy(d_x, x, N, cudaMemcpyHostToDevice);
cudaMemcpy(d_y, y, N, cudaMemcpyHostToDevice);

// Perform SAXPY on 1M elements
saxpy<<<4096,256>>>(N, 2.0, d_x, d_y);

cudaMemcpy(y, d_y, N, cudaMemcpyDeviceToHost);

```

<http://developer.nvidia.com/cuda-toolkit>

## 4

# SAXPY: THRUST C++ TEMPLATE LIBRARY

## Serial C++ Code (with STL and Boost)

```
int N = 1<<20;
std::vector<float> x(N), y(N);

...

// Perform SAXPY on 1M elements
std::transform(x.begin(), x.end(),
               y.begin(), y.end(),
               2.0f * _1 + _2);
```

[www.boost.org/libs/lambda](http://www.boost.org/libs/lambda)

## Parallel C++ Code

```
int N = 1<<20;
thrust::host_vector<float> x(N), y(N);

...

thrust::device_vector<float> d_x = x;
thrust::device_vector<float> d_y = y;

// Perform SAXPY on 1M elements
thrust::transform(d_x.begin(), d_x.end(),
                  d_y.begin(), d_y.begin(),
                  2.0f * _1 + _2);
```

<http://thrust.github.com>

## 5

## SAXPY: CUDA FORTRAN

## Standard Fortran

```

module mymodule contains
  subroutine saxpy(n, a, x, y)
    real :: x(:), y(:), a
    integer :: n, i
    do i=1,n
      y(i) = a*x(i)+y(i)
    enddo
  end subroutine saxpy
end module mymodule

program main
  use mymodule
  real :: x(2**20), y(2**20)
  x = 1.0, y = 2.0

  ! Perform SAXPY on 1M elements
  call saxpy(2**20, 2.0, x, y)

end program main

```

## Parallel Fortran

```

module mymodule contains
  attributes(global) subroutine saxpy(n, a, x, y)
    real :: x(:), y(:), a
    integer :: n, i
    attributes(value) :: a, n
    i = threadIdx%x+(blockIdx%x-1)*blockDim%x
    if (i<=n) y(i) = a*x(i)+y(i)
  end subroutine saxpy
end module mymodule

program main
  use cudafor; use mymodule
  real, device :: x_d(2**20), y_d(2**20)
  x_d = 1.0, y_d = 2.0

  ! Perform SAXPY on 1M elements
  call saxpy<<<4096,256>>>(2**20, 2.0, x_d, y_d)

end program main

```



## 6

# SAXPY: PYTHON

## Standard Python

```
import numpy as np

def saxpy(a, x, y):
    return [a * xi + yi
            for xi, yi in zip(x, y)]

x = np.arange(2**20, dtype=np.float32)
y = np.arange(2**20, dtype=np.float32)

cpu_result = saxpy(2.0, x, y)
```

<http://numpy.scipy.org>

## Numba: Parallel Python

```
import numpy as np
from numba import vectorize

@vectorize(['float32(float32, float32,
float32)'], target='cuda')
def saxpy(a, x, y):
    return a * x + y

N = 1048576

# Initialize arrays
A = np.ones(N, dtype=np.float32)
B = np.ones(A.shape, dtype=A.dtype)
C = np.empty_like(A, dtype=A.dtype)

# Add arrays onGPU
C = saxpy(2.0, A, B)
```

<https://numba.pydata.org>

# ENABLING ENDLESS WAYS TO SAXPY

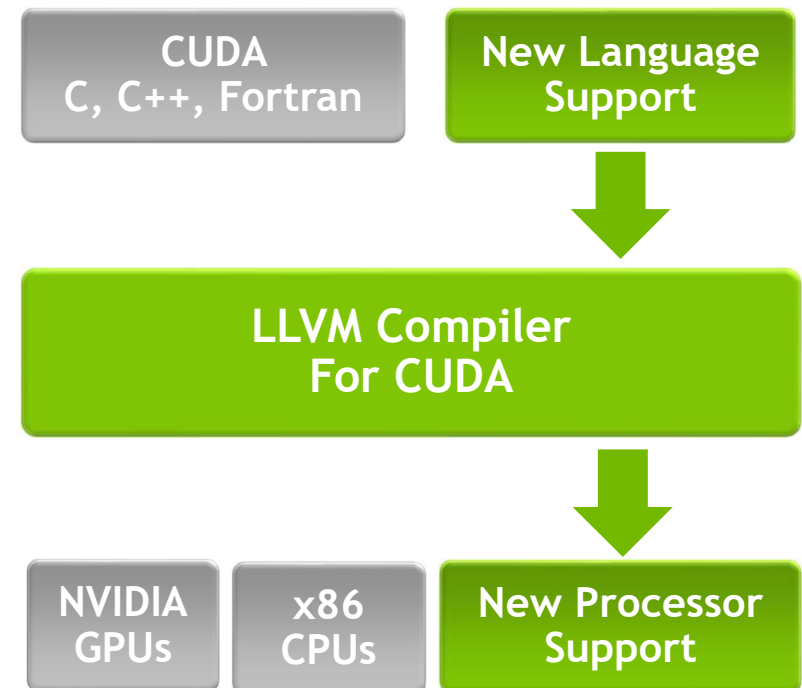
Developers want to build front-ends for:

- Java, Python, R, DSLs

Target other processors like:

- ARM, FPGA, GPUs, x86

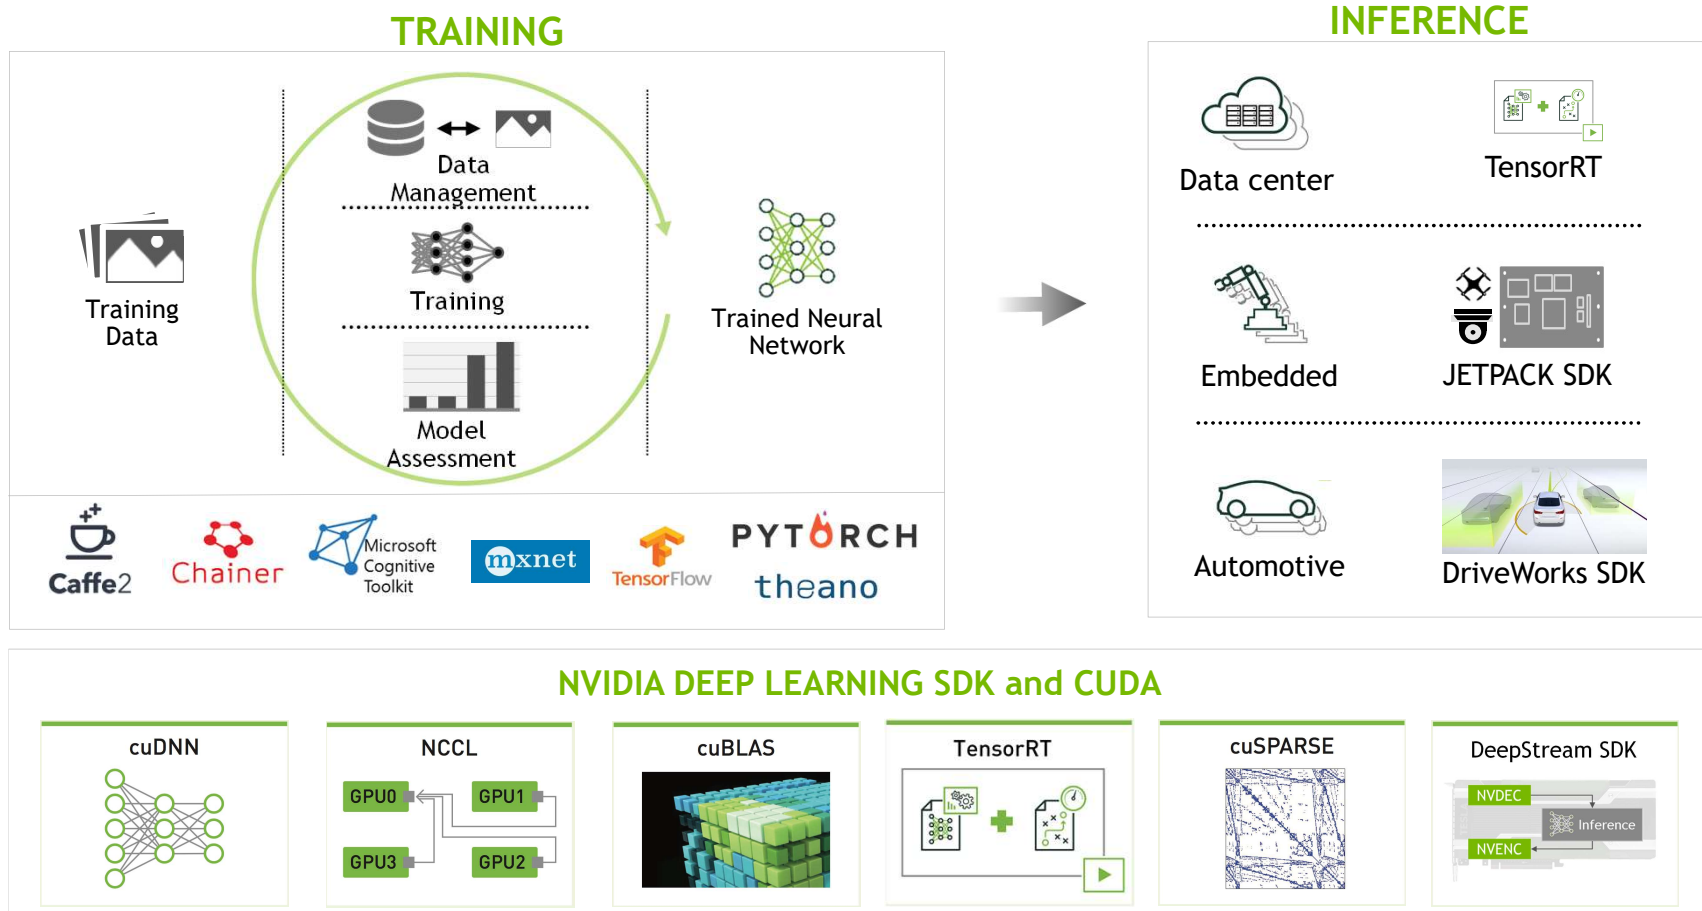
**CUDA Compiler Contributed  
to Open Source LLVM**



The background is a dark, almost black, field filled with a complex network of thin, glowing green lines. These lines intersect at various points, creating a web-like structure. At several of these intersection points, there are small, bright green circular dots. Additionally, there are a few larger, fainter blue circular shapes scattered across the background, giving it a sense of depth and a futuristic or technological feel.

# **GPU COMPUTING PLATFORM FOR DL/ML**

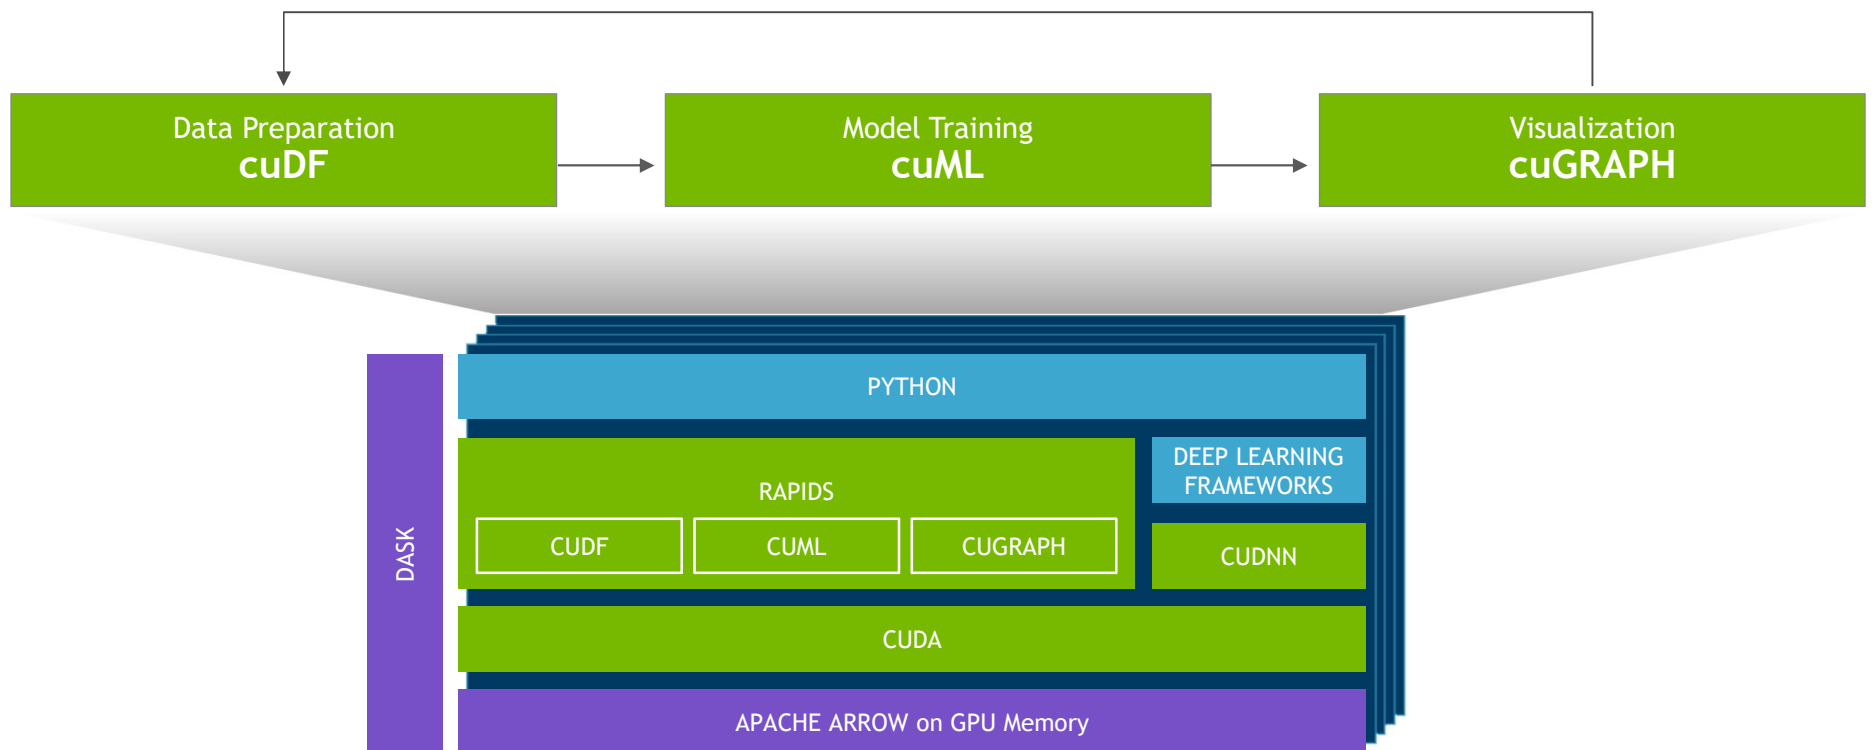
# NVIDIA DEEP LEARNING SOFTWARE STACK



[developer.nvidia.com/deep-learning-software](https://developer.nvidia.com/deep-learning-software)

# RAPIDS – OPEN GPU DATA SCIENCE

Software Stack Python





# WHY RAPIDS

## World's Fastest Machine Learning



### H2O.ai Machine Learning – Generalized Linear Modeling



U.S. Census dataset (predict Income): 45k rows, 10k cols  
Parameters: 5-fold cross-validation,  $\alpha = \{\frac{i}{7}, i = 0 \dots 7\}$ , full  $\lambda$ -search



An abstract network diagram with green nodes and lines on a dark background. The nodes are represented by small, glowing green circles of varying sizes, and the lines are thin, green, semi-transparent lines connecting the nodes. The lines crisscross the frame, creating a complex web of connections. The background is a deep black with some subtle, larger, out-of-focus green and blue light spots.

**NGC FOR EFFICIENCY**

# CHALLENGES UTILIZING AI & HPC SOFTWARE

## Installation




Complex, time consuming, and error-prone

## Optimization



Requires expertise to optimize framework performance

## Maintenance



IT can't keep up with frequent software upgrades

## Productivity



Users limited to older features and lower performance

# NGC

## The GPU-Optimized Software Hub



**Simplify Deployments with  
Performance-optimized Containers**



**Innovate Faster with Ready-to-Use  
Solutions**



**Deploy Anywhere**



<https://www.nvidia.com/gpu-cloud/>

# SUMMARY

- Full Stack Optimization is key to performance
- Multiple choices for programming on GPU
- One is not an alternative to other. They co-exist
- Universal hardware with Software stack is key to GPU computing



Thank You